



GUROBI OPTIMIZATION

**The Latest Advances in
the Gurobi Optimizer**

Recent Improvements

- ▶ Performance improvements
- ▶ New approaches for handling non-linear objective functions
- ▶ New remote computing capabilities
- ▶ New rich examples

New Rich Examples

Demo

- ▶ Rich example demo...


Piecewise-Linear Objective Functions

Motivation

- ▶ Non-linear objective functions – conventional wisdom:
 - Best way to model them is typically a piecewise-linear approximation
- ▶ Traditional PWL approach
 - One optimization variable for each piece
 - Increased accuracy \rightarrow more pieces \rightarrow rapid growth in problem size
- ▶ Objective of our “new” approach:
 - Extend simplex algorithm (primal and dual)
 - Handle PWL terms directly (no extra variables)
- ▶ New API to make it easy to express PWL functions


Convex PWL Objective Function

- ▶ Continuous
- ▶ Slopes are non-decreasing



Non-convex PWL Function

- ▶ Continuous non-convex function



PWL Simplex

- ▶ Extend both primal and dual simplex to support separable convex PWL objective
 - Primal PWL simplex, Fourer and Marsten, 1992
 - Dual PWL simplex, no reference
- ▶ Main difference between standard and PWL simplex
 - Ratio test
 - Depends on efficient implementation of the median algorithm
- ▶ Dual PWL simplex ~3X faster than primal PWL simplex

Computational Test, QP

- ▶ PWL approximation of QP objective
- ▶ Quality of approximation depends on number of pieces:
 - 100 pieces: 2 digits of accuracy
 - 1000 pieces: 4 digits of accuracy
- ▶ Compare PWL simplex against simplex on one-variable-per-piece model (for QP models that require > 1 s):

Pieces	Speedup
100	2.0X
1000	6.1X

- ▶ Compare PWL simplex (w/ 1000 PWL pieces) against QP simplex:
 - PWL only 2.3X slower than specialized QP simplex algorithm

Gurobi PWL APIs

► Point representation


- List points that define PWL function
 - Using (x,y) coordinates
- C routines

```
GRBsetpwlobj(model, var,
              npoints, x, y)
```
- OO APIs, e.g. C++

```
Model.setPWLObj
```

► Attributes


- NumPWLObjVars
- PWLObjCvx



Remote Computing – Compute Server, Distributed Optimization, and Gurobi Cloud


Gurobi Compute Server

- ▶ Client computer uses Gurobi API to build model
- ▶ Client computer passes model data to server
- ▶ Server solves the model
- ▶ Result values returned to client computer




Use Case – Thin Client

- ▶ Gurobi Compute Server
 - Offload optimization from slow clients to fast servers
 - One customer example:
 - 4X improvement in MIP solution times
 - Two servers (and two Gurobi licenses) handle all optimization jobs for 50+ clients
 - No specialized implementation work required




Use Case – Distributed Computing

- ▶ One manager offloads computation to multiple workers
- ▶ Multiple machines cooperate to solve a single MIP model



Use Case – Cloud Computing


- ▶ Gurobi Cloud hosted on Amazon EC2
- ▶ Launch an instance (or 100)
- ▶ Offload computation to it over the Internet
- ▶ Pay by the hour



Distributed MIP

Distributed MIP Architecture

- ▶ Manager-worker paradigm
- ▶ Manager
 - Send model to all workers
 - Track dual bound and worker node counts
 - Rebalance search tree to put useful load on all workers
 - Distribute feasible solutions
- ▶ Workers
 - Solve MIP nodes
 - Report status and feasible solutions
- ▶ Deterministic synchronization



Distributed MIP Phases

- ▶ Racing ramp-up phase
 - Distributed concurrent MIP
 - Solve same problem individually on each worker, using different parameter settings
 - Stop when problem is solved or “enough” nodes are explored
 - Choose a “winner” – worker that made the most progress
- ▶ Main phase
 - Discard all worker trees except the winner's
 - Collect active nodes from winner, distribute them among now idle workers
 - Periodically synchronize to rebalance load

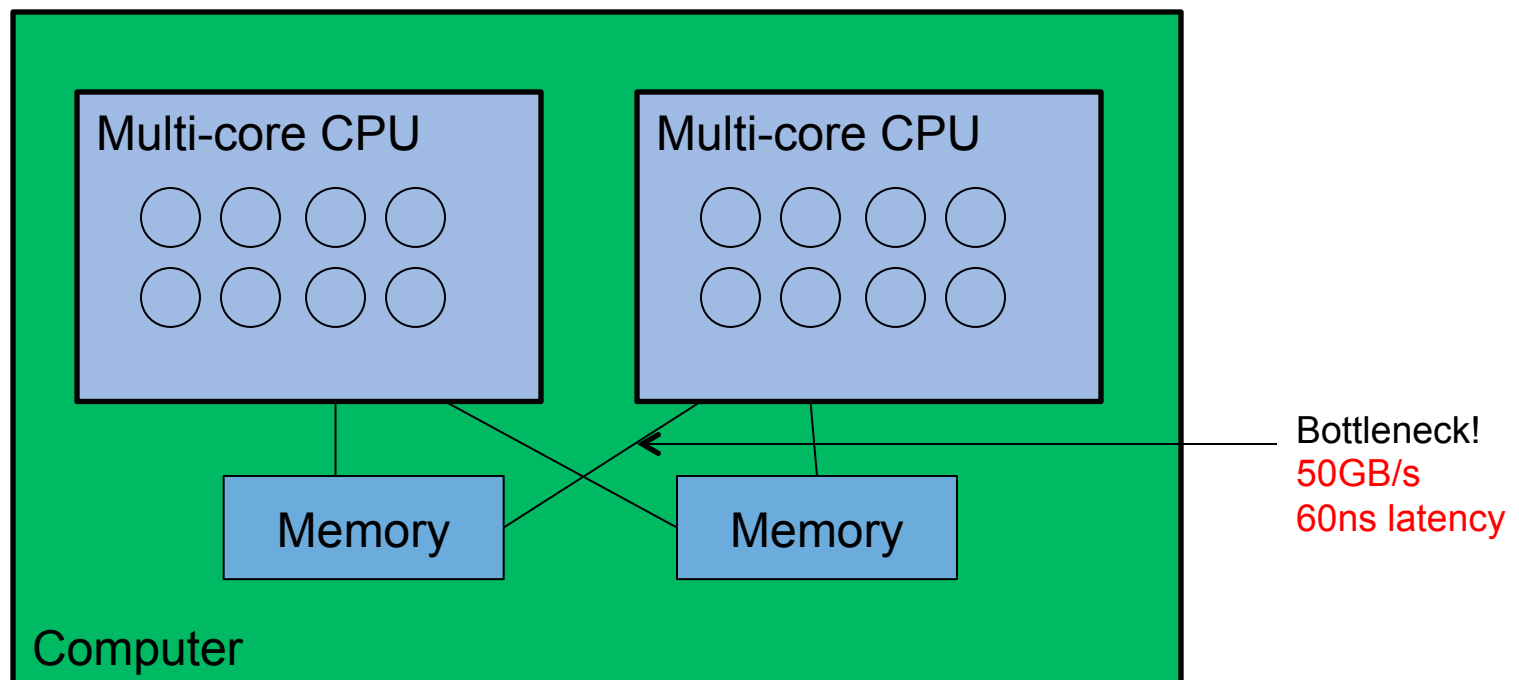
Distributed MIP Performance

16 Cores vs. 16 Cores

- ▶ Consider two different tests using 16 cores:
 - On a 16-core machine:
 - Run the standard parallel code on all 16 cores
 - On four 4-way machines:
 - Run the distributed code
- ▶ Which gives better results?

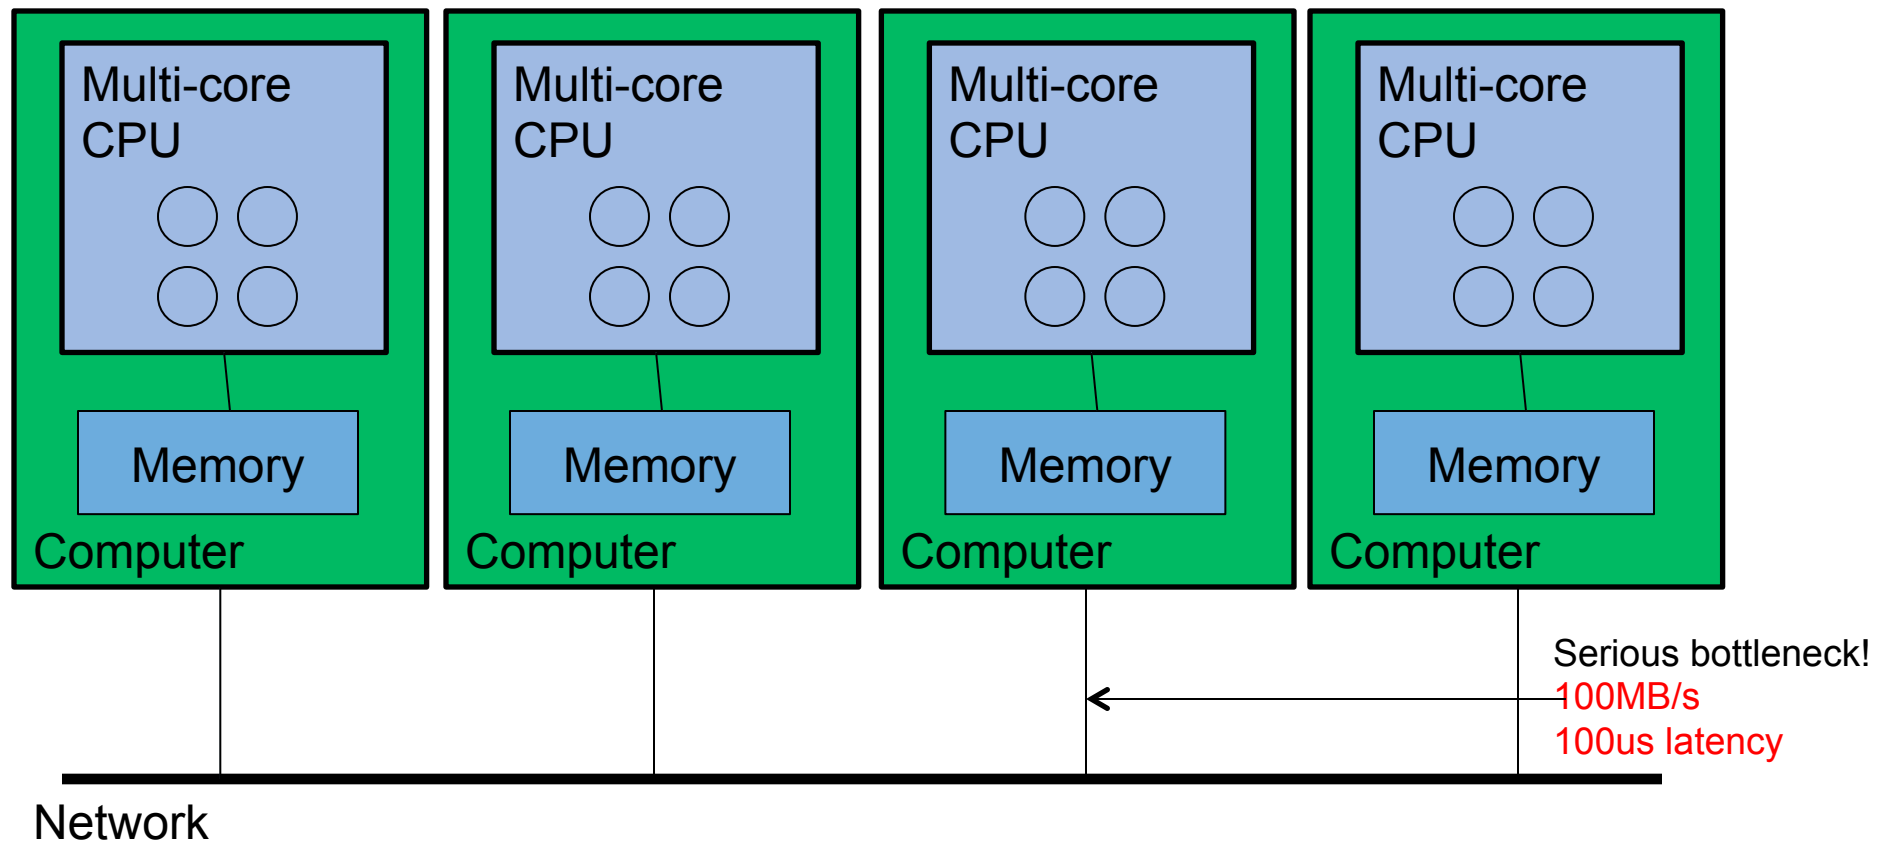
Parallel MIP on 1 Machine

- ▶ Use one 16-core machine:



Distributed MIP on 4 machines

- ▶ Use four 4-core machines



Performance Results

- ▶ Comparing one 16-core machine against four 4-core machines (MIPLIB 2010, baseline is single-machine, 4-core run)...

Config	>1s	>100s
One 16-core machine	1.57X	2.00X
Four 4-core machines	1.43X	2.09X

- ▶ Given a choice...
 - Comparable mean speedups
 - Other factors...
 - Cost: four 4-core machines are much cheaper
 - Admin: more work to admin 4 machines

Distributed Algorithms in 6.0

- ▶ MIPLIB 2010 benchmark set
 - Intel Xeon E3-1240v3 (4-core) CPU
 - Compare against 'standard' code on 1 machine

Machines	>1s			>100s		
	Wins	Losses	Speedup	Wins	Losses	Speedup
2	40	16	1.14X	20	7	1.27X
4	50	17	1.43X	25	2	2.09X
8	53	19	1.53X	25	2	2.87X
16	52	25	1.58X	25	3	3.15X

Some Big Wins

- ▶ Model *seymour*
 - Hard set covering model from MIPLIB 2010
 - 4944 constraints, 1372 (binary) variables, 33K non-zeroes

Machines	Nodes	Time (s)	Speedup
1	476,642	9,267s	–
16	1,314,062	1,015s	9.1X
32	1,321,048	633s	14.6X

Gurobi Distributed MIP

- ▶ Makes huge improvements in performance possible
- ▶ Mean performance improvements are significant but not huge
 - Much better than distributed concurrent
 - As effective as adding more cores to one box
- ▶ Effectively exploiting parallelism remains:
 - A difficult problem
 - A focus of ours

Gurobi Instant Cloud

Why Use The Cloud For Optimization?

- ▶ Pay just for what you use
 - Short-term projects
 - Occasional use
 - Meet a peak in demand
- ▶ No software or hardware to purchase or configure
- ▶ Get many fast computers quickly
 - Especially valuable for distributed optimization
- ▶ Increased robustness
 - Get computers in multiple locations
 - World-class management of computer center

Demo

- ▶ Gurobi Instant Cloud demo...

Pricing For Gurobi Instant Cloud

- ▶ Pay for your Cloud License plus the Machine Use

- ▶ Gurobi Cloud License

	Basic	Compute Server
Features	<ul style="list-style-type: none">• Single-use license• No distributed algorithms• No job queuing	<ul style="list-style-type: none">• Multiple concurrent uses• Distributed algorithms• Load balancing, job queuing
Cost	\$10/hr or \$1500/750 hrs/month	\$20/hr or \$3000/750 hrs/month
Minimum charge	30 minutes (\$5)	30 minutes (\$10)

- Distributed worker: no cost for Gurobi license
- No commitment – monthly price applies automatically at 150 hrs in calendar month

- ▶ Machine use

- Typically ~\$0.50/hr, rounded to next hour
- Varies based on machine type and location

Performance Improvements

Performance Improvements (6.0 vs 5.6)


Problem Class	> 1 sec	> 100 sec
	Speed-up	Speed-up
INTERNAL		
LP	1.06x	1.07x
MIP	1.17x	1.27x
MIQP	1.18x	1.20x
MIQCP	1.28x	2.07x
EXTERNAL		
Mittelmann "Easy"	1.15x	1.23x
Mittelmann Optimality Benchmark*	1.12x	1.09x

Notes:

1. 4 threads, 10000s time limit
2. * Average over 5 random seeds

MIP Performance Improvements

- ▶ Version-to-version improvements:
(Geometric mean speedup for models in our internal model set where at least one of the solvers takes more than 100s to solve)



Gurobi 6.5 Coming Soon

- ▶ New release coming soon
- ▶ Significant improvements

Thank You