

Model troubleshooting & performance tuning



GUROBI
OPTIMIZATION

Runtime issues

- ▶ Code fails to run
 - ▶ Unexpected result
 - ▶ Slow performance
-
- ▶ We will discuss
 - Specific symptoms of runtime issues
 - What causes them
 - How to troubleshoot them

Symptoms of code failing to run

- ▶ Blue screen (Windows)
- ▶ Segmentation fault (Linux, Mac)
- ▶ Computer is frozen
- ▶ No log or progress information
- ▶ No solution
- ▶ Error messages

Top causes of code failure

- ▶ Improperly written code
 - Failing to catch exceptions
 - Failing to test solution status
- ▶ License key issue
- ▶ Data errors
- ▶ Low on physical memory
- ▶ Hardware failure
- ▶ ...
- ▶ Gurobi bug

Diagnosing when code fails to run

- ▶ What exactly is failing?
- ▶ What do you need to do to make it fail?
- ▶ Can you reproduce it reliably or only randomly?
- ▶ Where in the code does it fail?
 - Model creation
 - Solving
 - Callbacks
 - Retrieving solution values

Troubleshooting code failure situations

- ▶ Determine when and where it occurs
- ▶ Catch exceptions
- ▶ Use system profiling tools
 - Ex: System Monitor, top, Activity Monitor, profiler in compiler, etc.
- ▶ Try model file independently from code
- ▶ Try different hardware or the cloud

Catching exceptions

- ▶ Easy to test for exceptions in OO interfaces:

```
try:
```

```
    m = read(sys.argv[1])
```

```
    m.optimize()
```

```
    for v in m.getVars():
```

```
        print v.VarName, v.X
```

```
except GurobiError as e:
```

```
    print "Error:", e
```

- ▶ With C, test the return code for every call to the Gurobi API
- ▶ Don't be sloppy – always test for exceptions!
 - Many support requests to Gurobi could be avoided by testing for exceptions and reviewing the exception values

Try model independently from code

- ▶ In your code, export the model as MPS file:

```
m.update();  
m.write("test.mps");
```
- Use MPS files, not LP files
 - Note: constant offset (ObjCon attribute) is *not* saved in MPS file
- ▶ Test model file separately:
 - Command-line utility

```
gurobi_cl test.mps
```
 - Python shell

```
m=read("test.mps")  
m.optimize()
```


Saving full state

- ▶ Use files to write and read parameters, warm starts
 - Non-default parameters: .PRM file
 - MIP start values: .MST file
 - Initial LP basis: .BAS file
- ▶ Export in your code:

```
m.update();  
m.write("test.mps");  
m.write("test.prm");  
m.write("test.mst");
```
- ▶ Test model file separately:
 - Command-line utility
`gurobi_cl inputfile=test.prm inputfile=test.mst test.mps`
 - Python shell

```
m=read("test.mps")  
m.read("test.prm")  
m.read("test.mst")  
m.optimize()
```

Symptoms of unexpected results

- ▶ Model is reported as infeasible
- ▶ A known solution is not feasible
- ▶ Solution appears to be suboptimal
- ▶ Solution violates some known constraints

Causes of unexpected results

- ▶ Modeling errors
- ▶ Unexpected data
- ▶ Numerical issues

Troubleshooting unexpected results

- ▶ Review model statistics
- ▶ Inspect LP file
- ▶ Test known solution
- ▶ Try infeasibility detection

Model statistics

- ▶ Does the model match your expectations for size, coefficients?
- ▶ `Model.printStats()` function in Python API gives a model overview

```
m=read('misc07')
m.printStats()
```

Statistics for model MISC07 :

Linear constraint matrix	: 212 Constrs, 260 Vars, 8619 NZs
Variable types	: 1 Continuous, 259 Integer
Matrix coefficient range	: [1, 700]
Objective coefficient range	: [1, 1]
Variable bound range	: [0, 1]
RHS coefficient range	: [1, 300]

- ▶ These stats are printed in the log, and they are available via the attribute interface

LP file

- ▶ LP file uses algebraic syntax
- ▶ With meaningful variable and constraint names, LP files can be a valuable debugging tool

Maximize

$x + y + 2z$

Subject To

c0: $x + 2y + 3z \leq 4$

c1: $x + y \geq 1$

Bounds

Binaries

x y z

End

Testing a known solution

MIP optimality

- ▶ Use solution as start values, then search for improvements
 - Use MIP start file (.MST) or Start variable attribute
- ▶ Python example

```
m=read("test.mps")
m.read("test.mst")
m.optimize()
```

Feasibility

- ▶ Set bounds to known solution
 - No automatic feature but easy to do with APIs
- ▶ Python example

```
m=read("test.mps")
x=m.getVarByName("x")
x.setAttr("LB", 1.0)
x.setAttr("UB", 1.0)
```

Infeasibility detection tools

IIS

- ▶ Finds a minimal set of constraints that conflict
- ▶ Attributes indicate what cannot be satisfied
- ▶ Useful for model development

FeasRelax

- ▶ Finds a solution with the minimum constraint violation
- ▶ Retrieve added artificial slack variables to find violations
- ▶ Useful for deployed application

Symptoms of slow performance



► It's slow!

Causes of slow performance

- ▶ Numerical issues
- ▶ Memory limits
- ▶ Unrealistic tolerance values
- ▶ A large or difficult model!

Troubleshooting slow performance

- ▶ Determine what cases cause the slow performance
- ▶ Determine what algorithmic part is the bottleneck
- ▶ Parameter tuning tool
- ▶ Performance guidelines
- ▶ Send test models to Gurobi

What cases cause slow performance

- ▶ What data cause slow performance?
 - Every case or selected conditions
- ▶ Is the slow performance predictable or seemingly random?
- ▶ Is it comparably slow on different computers?
 - Gurobi can provide temporary licenses for testing on other computers
 - Try the cloud

What algorithmic part is the bottleneck

- ▶ Model initialization and solution retrieval
 - Test MPS file using gurobi_cl; see if solution times are much faster
- ▶ Solve times
 - Presolve
 - Solving (initial LP)
 - At node 0 of MIP
 - Other nodes of MIP
 - Log shows time spent in presolve, LP relaxation, MIP root, nodes
- ▶ Use the logs to identify the bottleneck

grbtune: Parameter tuning tool

- ▶ Automatically finds sets of parameters that give good performance
- ▶ Two modes
 - Fastest time to optimality
 - Smallest MIP gap in a fixed amount of time
- ▶ Can run distributed across multiple computers
- ▶ An API is available
- ▶ **Best to tune with a representative set of models**

- ▶ Examples:

```
grbtune TuneTimeLimit=3600 modelA1.mps modelA2.mps modelA3.mps  
grbtune TuneTimeLimit=7200 TimeLimit=300 modelB1.mps modelB2.mps
```

Performance guidelines by category

- ▶ General issues
 - ▶ Continuous optimization
 - ▶ Integer optimization
-
- ▶ Let's examine these in detail

Performance guidelines by category

- ▶ General issues
- ▶ Continuous optimization
- ▶ Integer optimization

General issues

- ▶ Model initialization
- ▶ Presolve
- ▶ Numerical issues
- ▶ Memory

Model initialization

- ▶ Each matrix generator has its own pitfalls and best practices
 - Use iterators effectively for your API
- ▶ Look for: bottleneck via a code profiler
- ▶ With Gurobi OO interfaces, take advantage of lazy updates
 - Only call update function when necessary to reference new objects

Presolve

- ▶ Tradeoff: spend time up front with hope of simplifying model
- ▶ Look for: performance with different presolve parameters
- ▶ Primary control: **Presolve** parameter
 - Reduce if spending too much time up front
 - Increase to hope to get a simpler model
- ▶ Additional parameters for fine-grain control
 - **PrePasses**, **Aggregate**, **AggFill**

Numerical issues

- ▶ Wide range of model coefficients can ruin performance
 - Can affect continuous or integer models
- ▶ Look for: Unusual messages in solution log
 - Ex: Markowitz tolerance increased

Memory

- ▶ Insufficient memory can wreck performance
 - Virtual memory via disk is far slower than RAM
 - Parallel optimization requires more memory
- ▶ Look for: memory use via system monitor tools on computer
 - Ex: System Monitor, top, Activity Monitor
- ▶ Helpful parameters
 - Decrease [Threads](#)
 - Set [NodefileStart](#) to use disk to store MIP node info
- ▶ Memory is cheap; no need to skimp

Sources of bottlenecks

- ▶ General issues
- ▶ **Continuous optimization**
- ▶ Integer optimization

Continuous algorithms

- ▶ Dual simplex
- ▶ Primal simplex
- ▶ Barrier

- ▶ Concurrent (LP)
 - Use multiple algorithms at the same time on multiple processor cores
 - Multiple algorithms makes it very robust
 - Requires more memory

Concurrent

- ▶ Two concurrent modes
 - Fast
 - First algorithm that finishes wins
 - Occasionally, there is a near-tie
 - In this case, the solution can vary on multiple runs
 - Deterministic
 - Consistent, repeatable winner declaration
 - Usually much slower than fast mode
 - Winning method determined by internal metrics, not runtime

LP Performance

- ▶ Mean runtime ratios (quad-core Xeon E3-1240 v3):
 - ~210 models with runtime > 1 s
 - Concurrent: dual on 1 core, barrier on 3

	Geometric mean
Dual simplex	1.00
Primal simplex	2.11
Barrier	0.51
Concurrent	0.40
Deterministic concurrent	0.44

QP Performance

- ▶ Mean runtime ratios (quad-core i7-2600):
 - 14 models with runtime > 1 s

	Geometric mean
Dual simplex	1.00
Barrier	0.02

Typical defaults for continuous optimization

- ▶ LP Concurrent
- ▶ QP Barrier
- ▶ MIP root Dual simplex or concurrent, depending on model size
- ▶ MIP nodes Dual simplex

- ▶ Parameters used to select the algorithm
 - **Method**: continuous models and root of MIPs
 - **NodeMethod**: nodes of MIPs

LP – Example 1

- ▶ First run

```
gurobi> m.optimize()  
...  
Solved with dual simplex  
Solved in 11615 iterations and 3.72 seconds  
Optimal objective 2.382165864e+10  
gurobi> print m.getVars()[0].X  
351.0
```

- ▶ Second run of same model

```
gurobi> m.optimize()  
...  
Solved with barrier  
Solved in 53305 iterations and 3.70 seconds  
Optimal objective 2.382165864e+10  
gurobi> print m.getVars()[0].X  
0.0
```

LP – Example 1

- ▶ Default solver is concurrent
- ▶ Different optimal basis possible when dual and barrier runtimes are very close
- ▶ If this is an issue, use a deterministic method
 - Deterministic concurrent (can be much slower)
 - Parallel barrier
 - Simplex

LP – Example 2

- ▶ Solve model pds-100 on a 4-core i7-3770K
- ▶ Dual...
Solved in 112881 iterations and 15.36 seconds
- ▶ Barrier...
Solved in 78612 iterations and 40.41 seconds
- ▶ Concurrent...
Solved with dual simplex
Solved in 112881 iterations and 21.03 seconds
- ▶ “I thought you said that concurrent was the fastest option”

The Fastest Solver

- ▶ Concurrent fastest on average
- ▶ Concurrent *never* fastest for a specific model
 - If barrier wins:
 - Concurrent wasted one thread on dual
 - If dual wins:
 - Dual had to fight for resources with barrier
 - Most limiting resource: the cooling fan!

LP – Example 3

- ▶ Small LP solves quickly
- ▶ Larger LP solves exponentially slower
- ▶ “Is the disk activity light flashing on your PC?”
 - “Yes, why do you ask?”

Memory Usage

- ▶ Concurrent uses much more memory than dual
 - Concurrent invokes barrier
 - Barrier typically uses a lot more memory than dual
 - Concurrent runs multiple solvers at once
 - Each needs a copy of the model
- ▶ If memory is tight, use dual simplex

Notable continuous parameters

- ▶ **NormAdjust**
 - Select different simplex pricing norm variants
- ▶ **Crossover**
 - Select strategy used for transforming barrier solution to basic solution
 - No effect for QP or QCP models
- ▶ **CrossoverBasis**
 - Select strategy used to construct initial basis from barrier solution

Sources of bottlenecks

- ▶ General issues
- ▶ Continuous optimization
- ▶ Integer optimization

ID why your MIP is difficult

- ▶ Time to solve LP/QP relaxations?
- ▶ Moving the bound?
- ▶ Finding feasible solutions?

If relaxations are the bottleneck

- ▶ Use tuning methods for continuous optimization
 - Try different methods for root and nodes
 - Check for memory issues
 - Try different values of `NormAdjust` parameter

MIP – Example 1

- MIP log looks like this...

Nodes		Current Node			Objective Bounds			Work		
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time	
	0	0	-264137.12	0	75	-	-264137.12	-	0s	
H	0	0			-0.0000107	-264137.12	-	-	0s	
H	0	0			-162837.1173	-264137.12	62.2%	-	0s	
	0	0	-251769.54	0	50	-162837.12	-251769.54	54.6%	0s	
	0	0	-246602.68	0	83	-162837.12	-246602.68	51.4%	0s	
	0	0	-241768.49	0	29	-162837.12	-241768.49	48.5%	0s	
H	0	3			-180084.9071	-241768.49	34.3%	-	0s	
	0	3	-241768.49	0	29	-180084.91	-241768.49	34.3%	0s	
*	2343	1391		49	-181346.8006	-214776.82	18.4%	6.4	2s	
*	2440	1339		70	-183734.9437	-214277.60	16.6%	6.4	2s	
*	3103	1527		59	-183933.1807	-213309.24	16.0%	6.4	2s	
*	3710	1825		61	-184549.0628	-212449.79	15.1%	6.4	2s	
H	5876	3164			-188226.7818	-210900.89	12.0%	6.3	3s	
	9281	5152	-192568.56	62	50	-188226.78	-209633.03	11.4%	6.2	5s
	24765	12393	-194237.85	42	50	-188226.78	-205750.35	9.31%	6.4	10s
H	25267	11765			-190619.4788	-205716.49	7.92%	6.4	10s	
	39314	15223	cutoff	50	-190619.48	-203402.09	6.71%	6.4	15s	
	51427	17115	-194173.45	47	25	-190619.48	-201649.93	5.79%	6.4	20s
	67287	17611	cutoff	36	-190619.48	-199597.77	4.71%	6.4	25s	
	82144	15576	cutoff	42	-190619.48	-197694.67	3.71%	6.3	30s	
	96108	11245	cutoff	34	-190619.48	-195664.78	2.65%	6.3	35s	
	109391	4347	-192569.29	47	30	-190619.48	-193256.98	1.38%	6.1	40s

...

- What should you try?

MIP – Example 1

- ▶ Try changing the focus of the search...
 - MIPFocus=1: focus on finding feasible solutions
 - MIPFocus=2: focus on proving optimality

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
...									
	0	0	-209067.20	0	70	-162837.12	-209067.20	28.4%	1s
	0	5	-209067.20	0	70	-162837.12	-209067.20	28.4%	2s
H	79	66				-165743.5947	-208071.05	25.5%	3.7
H	1045	871				-167148.0384	-208070.60	24.5%	5.2
	1238	1014	-193529.37	37	55	-167148.04	-206993.76	23.8%	5.1
H	1647	1140				-168514.0578	-203113.35	20.5%	6.8
*	1787	994		61		-173659.7799	-203113.35	17.0%	7.0
*	2018	993		54		-175972.1389	-202292.18	15.0%	7.1
H	2692	615				-187397.9703	-200840.93	7.17%	6.9
	2829	688	-187728.90	32	47	-187397.97	-200681.75	7.09%	7.4
H	4228	578				-190619.4788	-198075.82	3.91%	7.2

MIP – Example 2

- ▶ MIP log looks like this...

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	164800.976	0 1096	1169768.62	164800.976	85.9%	-	186s
	0	0	165096.286	0 1607	1169768.62	165096.286	85.9%	-	526s
H	0	0			481589.35817	165096.286	65.7%	-	1065s
	0	0	165452.400	0 1705	481589.358	165452.400	65.6%	-	1419s
	0	0	165566.561	0 1774	481589.358	165566.561	65.6%	-	1783s
	0	0	165719.831	0 1778	481589.358	165719.831	65.6%	-	2368s
	0	0	165849.075	0 1924	481589.358	165849.075	65.6%	-	2819s
...									

- ▶ Issues:
 - Bound moving very slowly
 - “Stuck” at the root
- ▶ What should you try?

MIP – Example 2

- ▶ In extreme cases, try turning off cuts (set Cuts=0)...

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	164800.976	0	1108	1169768.62	164800.976	85.9%	- 225s
	0	0	164800.976	0	1004	1169768.62	164800.976	85.9%	- 249s
H	0	0			503233.56923	164800.976	67.3%	-	406s
	0	2	164800.976	0	805	503233.569	164800.976	67.3%	- 425s
	2	4	164888.973	2	1009	503233.569	164800.976	67.3%	4039 437s
	3	5	164800.976	2	889	503233.569	164800.976	67.3%	2861 440s
	5	7	164899.272	3	998	503233.569	164800.976	67.3%	2285 454s
	6	7	164998.929	4	1047	503233.569	164800.976	67.3%	2341 465s
	7	10	164837.100	4	1007	503233.569	164800.976	67.3%	2201 481s
	9	10	165308.611	5	1100	503233.569	164837.100	67.2%	2660 497s
	11	16	164837.100	5	1005	503233.569	164837.100	67.2%	2200 509s
	15	18	165265.546	6	1075	503233.569	164837.100	67.2%	2080 531s
	19	23	164837.100	6	1004	503233.569	164837.100	67.2%	1817 557s
	24	29	164864.772	7	1026	503233.569	164864.772	67.2%	1897 571s
H	28	30			435540.73217	164864.772	62.1%	1722	874s
H	29	32			371622.48999	164864.772	55.6%	1663	874s
...									

MIP – Example 2

- Change MIP strategies (set ImproveStartTime=3600)...

```
...      836      828 170012.416  113 1107 371622.490 164864.772  55.6%   901 3614s
```

Resetting heuristic parameters to focus on improving solution
(using Heuristics=0.5 and RINS=10)...

```
      913      905 170302.813  121 1073 371622.490 164864.772  55.6%   837 4349s
      917      906 170306.908  122 1073 371622.490 164864.772  55.6%   834 5120s
H   920      904                236882.29084 164864.772  30.4%   833 5120s
H   924      908                235325.82038 164864.772  29.9%   829 5120s
H   935      917                235325.82033 164864.772  29.9%   823 6310s
H   941      921                235325.81953 164864.772  29.9%   820 11402s
      945      930 173125.061  126 1055 235325.820 164864.772  29.9%   819 20667s
      949      936 172549.331  125 1077 235325.820 164864.772  29.9%   818 21413s
H   952      936                223946.40672 164864.772  26.4%   816 21413s
H   956      938                223050.75538 164864.772  26.1%   813 21413s
      957      942 172391.221  126 1052 223050.755 164864.772  26.1%   812 29280s
      961      940 173719.366  127 1048 223050.755 164864.772  26.1%   811 29963s
     1001      972 172401.160  130 1050 223050.755 164864.772  26.1%   794 30970s
H  1030      994                218867.25201 164864.772  24.7%   784 30970s
...
```

MIP – Example 3

- Solve progress slows, ‘top’ (or Task Manager) shows Gurobi not making good progress...

```
top - 14:27:11 up 22 days, 22:07, 3 users, load average: 4.15, 4.05, 3.99
Tasks: 73 total, 1 running, 71 sleeping, 1 stopped, 0 zombie
Cpu(s): 5.9%us, 0.5%sy, 0.0%ni, 32.4%id, 61.2%wa, 0.0%hi, 0.0%si, 0.0%st
Mem: 8193828k total, 8144716k used, 49112k free, 284k buffers
Swap: 19800072k total, 3337364k used, 16462708k free, 2108k cached
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
3414	rothberg	20	0	10.1g	7.4g	1636	D	23	95.3	657:27.82	gurobi_cl
207	root	15	-5	0	0	0	S	2	0.0	0:21.37	kswapd0
1	root	20	0	4020	168	168	S	0	0.0	0:01.22	init
2	root	15	-5	0	0	0	S	0	0.0	0:00.00	kthreadd
3	root	RT	-5	0	0	0	S	0	0.0	0:00.51	migration/0
4	root	15	-5	0	0	0	S	0	0.0	0:00.40	ksoftirqd/0
5	root	RT	-5	0	0	0	S	0	0.0	0:00.13	watchdog/0
6	root	RT	-5	0	0	0	S	0	0.0	0:00.54	migration/1

- What should you try?

MIP – Example 3

- ▶ Use node files
 - `NodefileStart` parameter
- ▶ Performance penalty typically less than 10%

MIP – Example 4

- Progress stalls, not short on memory...

Nodes		Current Node			Objective Bounds			Work	
Expl	Unexpl	Obj	Depth	IntInf	Incumbent	BestBd	Gap	It/Node	Time
	0	0	3817190.68	0	544	- 3817190.68	-	-	2s
H	0	0				1.555212e+09 3817190.68	100%	-	2s
H	0	0				1.540152e+09 3817190.68	100%	-	2s
H	0	0				1.496837e+09 3817190.68	100%	-	2s
H	0	0				1.639804e+08 3817190.68	97.7%	-	3s
...									
	5468	4240	3830021.31	435	96	8247474.49 3822445.39	53.7%	36.7	73s
H	5724	3251				7554727.8619 3822445.39	49.4%	35.7	73s
*	5724	3251		484		7554727.8619 3822445.39	49.4%	35.7	73s
H	6526	921				3830692.9887 3822445.39	0.22%	33.0	74s
	6526	921		535		3830692.9887 3822445.39	0.22%	33.0	74s
...									
	22945	6545	3824602.97	208	45	3824990.61 3824448.25	0.01%	27.5	175s
*	23183	6450		348		3824966.4316 3824448.25	0.01%	27.3	175s
	27795	10035	3824931.45	147	52	3824966.43 3824448.56	0.01%	25.6	180s
...									
	3975393	2423369	cutoff	123		3824938.94 3824541.29	0.01%	23.0	7200s
...									

- What should you try?

MIP – Example 4

- ▶ Adjust your termination criteria
- ▶ Model data often have estimation errors
- ▶ Default MIPGap (0.01%) is overkill for many models

Parameter pitfalls

- ▶ Don't over-tune parameters
 - Default values are carefully selected, based on thousands of models
 - Avoid setting parameters unless they make a big improvement across multiple test models
- ▶ Don't assume parameters that were effective for another solver are ideal for Gurobi

More resources

- ▶ Sections in the Reference Manual
 - Logging
 - Parameter Guidelines
 - Parameter Tuning
- ▶ Gurobi support: support@gurobi.com